

Déploiement live d'un cluster K3s sur Linux

Présentation rapide du contexte

Ce déploiement vise à offrir une introduction concrète à Kubernetes via l'installation et l'utilisation d'un cluster léger K3s sous Linux. K3s est une version allégée de Kubernetes, adaptée à des environnements de test ou à des ressources limitées, simplifiant ainsi les premiers pas dans l'écosystème Kubernetes.

```
oot@debian:~# curl -sL https://get.k3s.io | sh -
INFO] Finding release for channel stable
INFO] Using v1.32.5+k3s1 as release
INFO] Downloading hash https://github.com/k3s-io/k3s/releases/download/v1.32.5+k3s1/sha256sum-amd64.txt
INFO] Downloading binary https://github.com/k3s-io/k3s/releases/download/v1.32.5+k3s1/k3s
INFO] Verifying binary download
INFO] Installing k3s to /usr/local/bin/k3s
INFO] Skipping installation of SELinux RPM
INFO] Creating /usr/local/bin/kubectl symlink to k3s
INFO] Creating /usr/local/bin/crictl symlink to k3s
INFO] Creating /usr/local/bin/ctr symlink to k3s
INFO] Creating killall script /usr/local/bin/k3s-killall.sh
INFO] Creating uninstall script /usr/local/bin/k3s-uninstall.sh
INFO] env: Creating environment file /etc/systemd/system/k3s.service.env
INFO] systemd: Creating service file /etc/systemd/system/k3s.service
INFO] systemd: Enabling k3s unit
created symlink /etc/systemd/system/multi-user.target.wants/k3s.service → /etc/systemd/system/k3s.service.
INFO] Host iptables-save/iptables-restore tools not found
INFO] Host ip6tables-save/ip6tables-restore tools not found
INFO] systemd: Starting k3s
```

Test

```
root@debian:/home/enzo# cp /etc/rancher/k3s/k3s.yaml ~/.kube/config
root@debian:/home/enzo# chown $(id -u):$(id -g) ~/.kube/config
root@debian:/home/enzo# export KUBECONFIG=~/.kube/config
root@debian:/home/enzo# kubectl get nodes
NAME      STATUS    ROLES                  AGE      VERSION
debian    Ready     control-plane,master   3m3s     v1.32.5+k3s1
root@debian:/home/enzo#
```

On peut voir que ça fonctionne et qu'un node remonte

Intégration de Helm

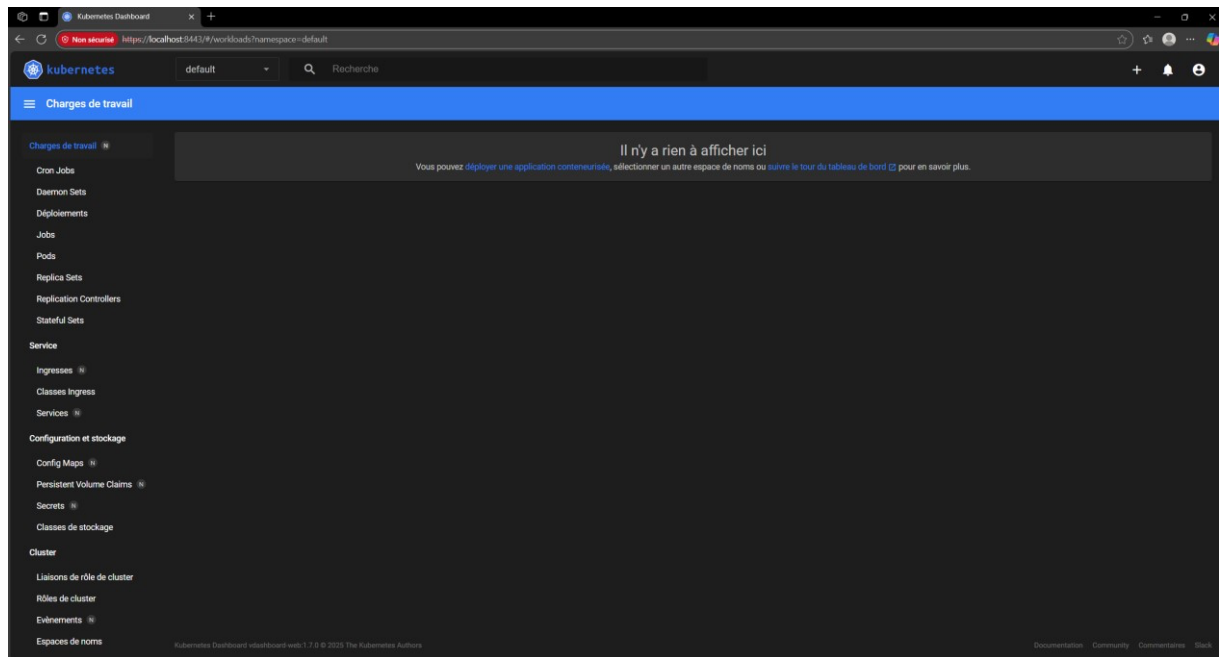
L'installation rapide de Helm a facilité le déploiement de charts, ensembles structurés de ressources Kubernetes.

```
curl -fsSL -o get_helm.sh
https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3
chmod +x get_helm.sh
./get_helm.sh
```

helm version

Pour tester que l'installation c'est faite correctement

```
root@debian:/home/enzo# kubectl port-forward -n kubernetes-dashboard svc/kubernetes-dashboard-kong-proxy 8443:443
Forwarding from 127.0.0.1:8443 -> 8443
Forwarding from [::1]:8443 -> 8443
Handling connection for 8443
Handling connection for 8443
Handling connection for 8443
```



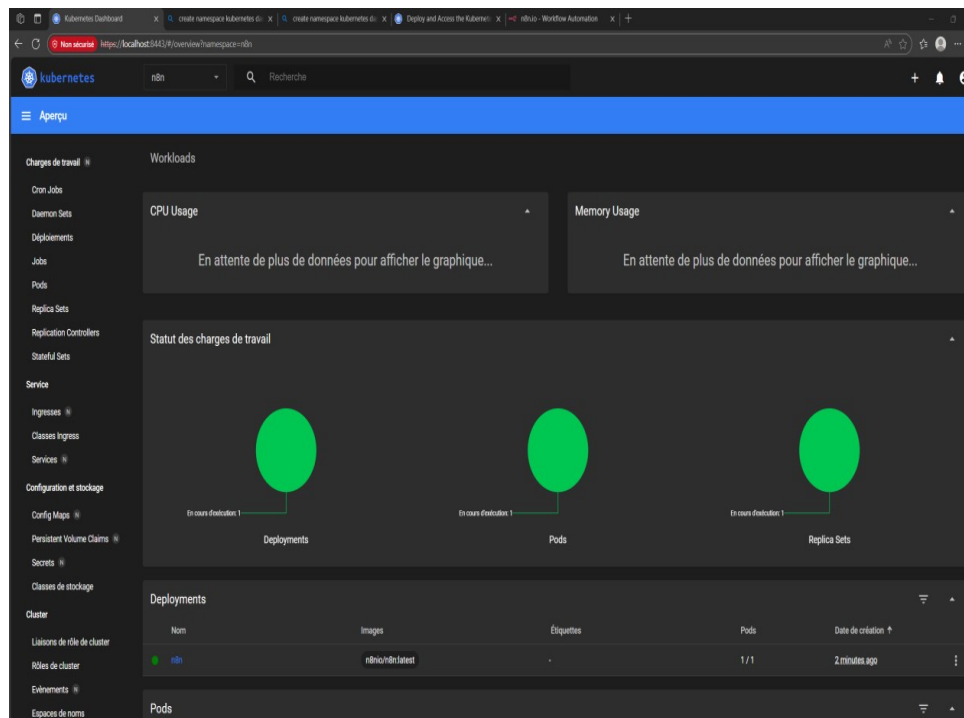
Ensuite j'ai créé le namespace N8N depuis le Dashboard en cliquant sur le +

J'ai cliqué à nouveau sur le + pour y déposer mon YML afin de déployer un node n8n

```

yaml
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: n8n-pvc
  namespace: n8n
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
      storageClassName: local-path
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: n8n
  namespace: n8n
spec:
  replicas: 1
  selector:
    matchLabels:
      app: n8n
  template:
    metadata:
      labels:
        app: n8n
    spec:
      containers:
        - name: n8n
          image: n8nio/n8n:latest
          ports:
            - containerPort: 5678
          volumeMounts:
            - name: n8n-storage
              mountPath: /home/node/.n8n
          env:
            - name: DB_TYPE
              value: "sqlite"
            - name: N8N_BASIC_AUTH_ACTIVE
              value: "true"
            - name: N8N_BASIC_AUTH_USER
              value: "admin"
            - name: N8N_BASIC_AUTH_PASSWORD
              value: "yourpassword"
          volumes:
            - name: n8n-storage
              persistentVolumeClaim:
                claimName: n8n-pvc
---
apiVersion: v1
kind: Service
metadata:
  name: n8n
  namespace: n8n
spec:
  type: NodePort
  ports:
    - port: 5678
      targetPort: 5678
      nodePort: 30001
  selector:
    app: n8n

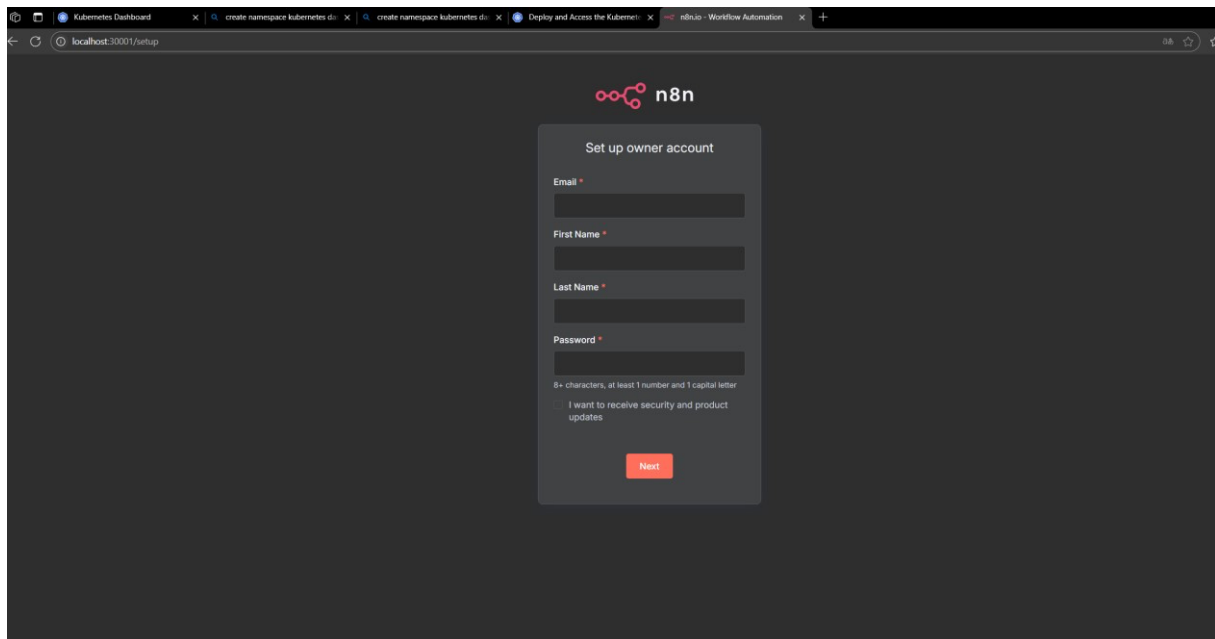
```



Ci-dessus un aperçu du déploiement réussi

On peut voir que ça fonctionne après avoir redirigé le port 3001 de ma VM sur mon pc

```
PS C:\Users\Enzo> ssh -L 30001:localhost:30001 enzo@192.168.65.171
```



Problèmes Rencontrés

A la connexion sur le Dashboard j'avais une erreur qui était du a un manque de ram comme j'ai pu le voir ci-dessous dans les logs

```
memory: 488Mi
Requests:
cpu: 100m
memory: 200Mi
Environment:
  CSRF_KEY: <set to the key 'private.key' in secret 'kubernetes-dashboard-csrf'> Optional: false
  GOMAXPROCS: 1 (limits.cpu)
  GOMEMLIMIT: 419430400 (limits.memory)
Mounts:
  /tmp from tmp-volume (rw)
  /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-jfw4t (ro)
Conditions:
  Type              Status
  PodScheduled      False
Volumes:
  tmp-volume:
    Type:            EmptyDir (a temporary directory that shares a pod's lifetime)
    Medium:
    SizeLimit:        <unset>
  kube-api-access-jfw4t:
    Type:            Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:    kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI:      true
QoS Class:           Burstable
Node-Selectors:       <none>
Tolerations:         node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                     node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type      Reason              Age           From          Message
  ---      -
  Warning   FailedScheduling    38m          default-scheduler  0/1 nodes are available: 1 Insufficient memory. preemption: 0/1 nodes are available: 1 No preemption victims found for incoming pod.
  Warning   FailedScheduling    12m (x5 over 32m) default-scheduler  0/1 nodes are available: 1 Insufficient memory. preemption: 0/1 nodes are available: 1 No preemption victims found for incoming pod.
root@debian:/home/enzo#
```